

CS143 Notes: Database Integrity

Book Chapters

Chapter 6.1-4

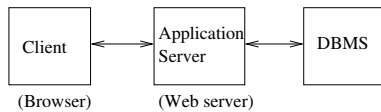
Things to Learn

- Key constraints
- Referential integrity (Foreign key constraints)
- CHECK constraints
- SQL assertion
- SQL trigger (part of SQL99)

What are integrity constraints?

- An example database with invalid entries (Show the example)
- A statement about what a valid database should look like
 - As a human being, we understand what is a “valid” database
 - The system needs an explicit specification of the semantics/rules
- Arbitrary predicate pertaining to the database (in principle)
 - In practice, only the ones that are easy to enforce
- If a SQL statement violates IC, the statement is aborted and generates an error
- **Q:** What rules/constraints can you find from the example?

- Database constraints checks the rules in the DB (Three tier diagram)



- **Q:** Why do we check these rules in DB, not in application?
Checking them at application/Web browser can be cheaper

Data validity enforcement in RDBMS

- 3 ways to enforce data validity in RDBMS
 - Domain: GPA is real
 - Constraints: Gives error. Abort statement
 - * Key
 - * Referential Integrity
 - * CHECK constraint
 - * SQL assertions
 - Trigger: Event-Condition-Action rule. If a certain event happens, invoke an action to handle it

Key Constraints

- A set of attributes should be unique in a table
- Course(dept, cnum, sec, unit, instructor, title)
Course(dept, cnum, sec, unit, instructor, title)
Course(dept, cnum, sec, unit, instructor, title)
 - CREATE TABLE Course (
 - dept CHAR(2) NOT NULL,
 - cnum INTEGER NOT NULL,
 - sec INTEGER NOT NULL,
 - unit INTEGER,
 - instructor VARCHAR(30),
 - title VARCHAR(30),
 - PRIMARY KEY(dept, cnum, sec),
 - UNIQUE(dept, cnum, instructor),
 - UNIQUE(dept, sec, title)

- One primary key per table
- Unique for other keys
- Primary key, unique are enforced through index (more discussion later)


Referential Integrity Constraints

• Example:

- If an sid appears in Enroll, it should also appear in Student
- If an (dept, cnum, sec) appears in Enroll, it should also appear in Class
 - * **Q:** Is the reverse true?

• Terminology

- (Two table diagram: E.A references C.A)



- E.A **references** C.A
- E.A: referencing attribute or **foreign key**
- C.A: referenced attribute
- **Referential integrity** means that referenced value always exists
 - * **foreign key can be NULL. When a foreign key is NULL, no constraint checking**

• Referential Integrity in SQL

– Example:


```
CREATE TABLE Enroll (
    sid INTEGER REFERENCES Student(sid),
    dept CHAR(2),
    cnum INTEGER,
    sec INTEGER,
    FOREIGN KEY (dept, cnum, sec) REFERENCES Class(dept, cnum, sec) )
```

– Notes:

- * Referenced attributes must be PRIMARY KEY or UNIQUE
- * Referenced attributes may be omitted if they are the same name with referencing attributes
 - e.g., sid INT REFERENCES Student
- * One attribute foreign key may be defined directly

• Referential Integrity Violation

- **Q:** When is the RI violated (two table diagram)?



e.g., do we have to worry if a tuple is deleted from E?

- RI violation from E (insert to E or update to E.A) is not allowed
 - * System rejects the statement
 - * Always insert/update C first.

- **Q:** If a tuple in C is updated/deleted, what can we do to avoid RI violation?

- **Comments:** Referential integrity is the only SQL constraint that can “fix itself”
 - * Other constraints simply abort and report error

- **Q:** Why should the referenced attributes be unique?

- Self referencing table

– **Example:**

A	B
1	NULL
2	1
3	2
4	3
5	4


```
CREATE TABLE R (  
    A INTEGER PRIMARY KEY,  
    B INTEGER REFERENCES R(A)  
    ON DELETE CASCADE )
```

– **Comments:**

- * A table references itself: self-referencing table
- * **Q:** What will happen if we delete (1,NULL)?

• **Circular constraints**

- **Example:** `ChickenFrom(cid, eid): eid became cid,`
`EggFrom(eid, cid): eid is born of cid`
 $(\text{Chicken.eid} \subset \text{Egg.eid}, \text{Egg.cid} \subset \text{Chicken.cid})$ (diagram)



- **Q:** Can we insert any tuple to Chicken? or to Egg? How can we fix it?

CHECK constraint

- Constraints attached to a table
- CHECK clause in table definition
- **Example:** $0 \leq GPA \leq 4.0$

```
CREATE TABLE Student(
    sid int,
    ...
    GPA real CHECK(0 <= GPA and GPA <= 4.0),
    ...)
```

- **Example:** `cnum < 600 AND unit < 10`

```

CREATE TABLE Enroll(
    dept CHAR(2),
    cnum INT,
    unit INT,
    title VARCHAR(50),
    CHECK (cnum < 600 AND unit < 10) )

```

- Constraint is checked whenever the tuple updated.
- In SQL92, conditions can be complex, e.g., with subqueries
- **Q:** The units of all CS classes are above 3 for Class(dept, cnum, unit, title)?

- **Q:** Students whose GPA is below 2.0 cannot take CS classes?

- **Q:** Can we express referential integrity constraint, e.g., Enroll.sid \subset Student.sid, using CHECK?

SQL Assertions

- Constraints on entire set of relations or database (not on the update of a single tuple).
- CREATE ASSERTION <name> CHECK (<condition>)
- **Example:** Average GPA > 3.0


```

      -- CREATE ASSERTION HighGPA
      CHECK(3.0 < (SELECT AVG(GPA) FROM Student))

```

- **Q:** What operations could violate the assertion HighGPA?

- The system figures out which operations may violate the constraints and monitors such operations.
- **Example:** A student can take at most 20 units in total

- **Q:** Can we express referential integrity using assertions?

Triggers

- After SQL92, people realized that it is very difficult to figure out when to check assertion automatically. Check almost every modification. Too expensive. Nobody implements it.
- For SQL99, people came up with a “easier” and “more procedural” version of assertion that explicitly lists the events to monitor and the action to take.

Trigger

- Event-Condition-Action rule (or ECA rule)
 - We explicitly specify what events to monitor, what condition to check and what action to take if the condition is met.
- **Query 1:** All new students have to take CS143 (For every insertion to Student, add the corresponding tuple to Enroll.)

Q: What if we insert thousands of student tuples in one insertion? Execute the trigger thousand times? Can we execute it once for all new tuples?

- **Query 2:** If a student GPA is updated to less than 2.0, revert back to the old GPS.

- Trigger general syntax: Event-Condition-Action rule (or ECA rule)
 - CREATE TRIGGER <name>
 <event>


```

<referencing clause>// optional
WHEN (<condition>) // optional
<action>

```

– <event>

```

* BEFORE | AFTER INSERT ON R
* BEFORE | AFTER DELETE ON R
* BEFORE | AFTER UPDATE [OF A1, A2, ..., An] ON R

```

– <action>

```

* Any SQL statement. Multiple statements should be enclosed with BEGIN ... END
  and be separated by ;

```

– <referencing clause>

```

* REFERENCING OLD|NEW TABLE|ROW AS <var>, ...
* FOR EACH ROW: row-level trigger
* FOR EACH STATEMENT (default): statement-level trigger

```

- **Query 3:** How to implement HighGPA assertion using triggers?

```

CREATE ASSERTION HighGPA
CHECK(3.0 < (SELECT AVG(GPA) FROM Student))

```

- **Q:** For, $R(A)$, after inserting (1), what will happen?

```

CREATE TRIGGER Recursion
AFTER INSERT ON R
BEGIN INSERT INTO R VALUES (1); END

```

What is supported in MySQL

- Key constraint
- Under InnoDB, most referential integrity except “ON DELETE/UPDATE SET DEFAULT”
- No CHECK constraints or assertions
- No assertions
- Limited trigger: does not allow updating the table that caused the trigger event
 - Generates error and rejects the statement that caused the event

Things to Remember

Constraints and Trigger

- Key constraint: PRIMARY KEY, UNIQUE
- Referential Integrity
 - Referencing attribute (foreign key), referenced attribute
 - * Referenced attribute should be PRIMARY KEY or UNIQUE
 - Violation at referencing attribute not allowed
 - Violation at referenced attribute can be fixed automatically
 - * ON DELETE/UPDATE SET NULL/SET DEFAULT/CASCADE
- Tuple-based CHECK constraint
- SQL assertions
- Trigger