



CS143: Disks and Files

1

System Architecture




2




- Capacity keeps increasing, but speed does not

3




4



5

Structure of a Platter



- Track, cylinder, sector (=block, page)


6

A Typical Disk

- Platter diameter: 1-5 in
- Cylinders: 100 – 2000
- Platters: 1 – 20
- Sectors per track: 200 – 500
- Sector size: 512 – 50K
- Rotation speed: 1000 – 15000 rpm
- Overall capacity: 1G – 300GB
- Q: 2 platters, 2 surfaces/platter, 500 tracks/surface, 200 sect/track, 1KB/sector. What is the overall capacity?

7

Access Time



- Q: How long does it take to read a page of a disk to memory?
- Q: What needs to be done to read a page?

8


Access Time

- Access time =
(seek time) + (rotational delay) +
(transfer time)

9

Seek Time


- Time to move a disk head between tracks



- Track to track ~ 1ms
- Average ~ 10 ms
- Full stroke ~ 20 ms

10


Rotational Delay



- Typical disk:
– 1000 rpm – 15000 rpm
- Q: For 6000 RPM, average rotational delay?

11

Transfer Rate



6000 RPM, 400 sectors/track, 512B/sector

- Q: How long to read one block?
- Q: What is the transfer rate (bytes/sec)?

12


(Burst) Transfer Rate

- (Burst) Transfer rate =
 $(\text{RPM} / 60) * (\text{sectors/track}) * (\text{bytes/sector})$

13

Sequential vs. Random I/O

- Q: How long to read 3 sequential blocks?




- ❑ 6000 RPM
- ❑ 200 sectors/track
- ❑ Assume the head is above the first block

14

Sequential vs. Random I/O

- Q: How long to read 3 random blocks?



- ❑ 6000 RPM
- ❑ 200 sectors/track
- ❑ 10ms seek time
- ❑ Assume the head is above the first block

15

Important to remember

- Random I/O: VERY expensive
 - Compared to sequential I/O
- Avoid random I/O as much as we can


16

Data Modification


- Byte-level modification not allowed
 - Can be modified by blocks
- Q: How can we modify only a part of a block?

17

Abstraction by OS



(head, cylinder, sector)



- Sequential blocks
 - No need to worry about head, cylinder, sector
- Access to non-adjacent blocks
 - Random I/O
- Access to adjacent blocks
 - Sequential I/O

18

Buffers, Buffer pool

- Temporary main-memory “cache” for disk blocks
 - Avoid future read
 - Hide disk latency
 - Most DBMS let users change buffer pool size

19

Reference

- Storage review disk guide
 - <http://www.storagereview.com/guide2000/ref/hdd/index.html>

20

Files: Main Problem

- How to store tables into disks?

Jane	CS	3.7
Susan	ME	1.8
June	EE	2.6
Tony	CS	3.1



21



Spanned vs Unspanned

- Q: 512Byte block. 80Byte tuple. How to store?



22

Spanned vs Unspanned

- Unspanned



- Spanned




- Q: Maximum space waste for unspanned?

23

Variable-Length Tuples

- How do we store them?



24

Reserved Space


- Reserve the maximum space for each tuple



- Q: Any problem?

25


Variable-Length Space



- Pack tuples tightly
- Q: How do we know the end of a record?
- Q: What to do for delete/update?
- Q: How can we "point to" a tuple?

26

Slotted Page



Q: How can we point to a tuple?

27


Long Tuples

- ProductReview(
pid INT,
reviewer VARCHAR(50),
date DATE,
rating INT,
comments VARCHAR(1000))
- Block size 512B
- How should we store it?

28

Long Tuples

- Spanning
- Splitting tuples



29

Sequential File

- Tuples are ordered by certain attribute(s) (search key)


Elaine	CS	3.7
James	ME	2.8
John	EE	1.8
Peter	EE	3.9
Susan	CS	1.0
Tony	EE	2.4

– Search key: Name

30

Sequencing Tuples

- Inserting a new tuple
 - Easy case



31


Sequencing Tuples

Two options

1) Rearrange

T1
T3
T6
T7
T8


2) Linked list



32

Sequencing Tuples


- Inserting a new tuple
 - Difficult case



33

Sequencing Tuples

- Overflow page



- Reserving free space to avoid overflow
 - PCTFREE in DBMS
- CREATE TABLE R(a int) PCTFREE 40

34

Sequencing Tuples Initially

- CREATE TABLE T ...;
INSERT INTO T
 (SELECT * ... ORDER BY key);
- Future insertions will gradually destroy the order
 - Periodic reordering may be necessary
- Other possibilities discussed later

35

Things to Remember

- Disk
 - Platter, track, cylinder, sector, block
 - Seek time, rotational delay, transfer time
 - Random I/O vs Sequential I/O
- Files
 - Spanned/unspanned tuples
 - Variable-length tuples (slotted page)
 - Long tuples
 - Sequential file and search key
 - Problems with insertion (overflow page)
 - PCTFREE

36