

CS143: Index

Book Chapters:
(4th) 12.1-3, 12.5-8
(5th) 12.1-3, 12.6-8, 12.10

1

Topics to Learn

- Important concepts
 - Dense index vs. sparse index
 - Primary index vs. secondary index
(= clustering index vs. non-clustering index)
 - Tree-based vs. hash-based index
- Tree-based index
 - Indexed sequential file
 - B+-tree
- Hash-based index
 - Static hashing
 - Extendible hashing

2

Basic Problem

- `SELECT *`
`FROM Student`
`WHERE sid = 40`

sid	name	GPA
20	Elaine	3.2
70	Peter	2.6
40	Susan	3.7

- How can we answer the query?

3

Random-Order File

- How do we find sid=40?

sid	name	GPA
20	Susan	3.5
60	James	1.7
70	Peter	2.6
40	Elaine	3.9
30	Christy	2.9

4

Sequential File

- Table sequenced by sid. Find sid=40?

sid	name	GPA
20	Susan	3.5
30	James	1.7
40	Peter	2.6
50	Elaine	3.9
60	Christy	2.9

5


Binary Search

- 100,000 records
- Q: How many blocks to read?
- Any better way?
 - In a library, how do we find a book?

6

Basic Idea


- Build an “index” on the table
 - An auxiliary structure to help us locate a record given a “key”



7

Dense, Primary Index

Dense Index Sequential File



- Primary index (clustering index)
 - Index on the search key
- Dense index
 - (key, pointer) pair for every record
- Find the key from index and follow pointer
 - Maybe through binary search
- Q: Why dense index?
 - Isn't binary search on the file the same?

8


Why Dense Index?

- Example
 - 10,000,000 records (900-bytes/rec)
 - 4-byte search key, 4-byte pointer
 - 4096-byte block. Unspanned tuples
- Q: How many blocks for table (how big)?
- Q: How many blocks for index (how big)?

9

Sparse, Primary Index

Sparse Index Sequential File




- Sparse index
 - (key, pointer) pair per every “block”
 - (key, pointer) pair points to the first record in the block
- Q: How can we find 60?

10

Multi-level index

Sparse 2nd level 1st level Sequential File




Q: Why multi-level index?

Q: Does dense, 2nd level index make sense?

11

Secondary (non-clustering) Index


Sequence field



- Secondary (non-clustering) index
 - When tuples in the table are not ordered by the index search key
 - Index on a non-search-key for sequential file
 - Unordered file
- Q: What index?
 - Does sparse index make sense?


12

Sparse and secondary index?



13

Secondary index



- First level is always dense
- Sparse from the second level

14


Important terms

- Dense index vs. sparse index
- Primary index vs. secondary index
 - Clustering index vs. non-clustering index
- Multi-level index
- Indexed sequential file
 - Sometimes called ISAM (indexed sequential access method)
- Search key (≠ primary key)

15

Insertion

Insert 35




Q: Do we need to update higher-level index?

16

Insertion

Insert 15 (overflow)




Q: Do we need to update higher-level index?

17

Insertion

Insert 15 (redistribute)




Q: Do we need to update higher-level index?

18

Potential performance problem

After many insertions...



19

Traditional Index (ISAM)

- Advantage
 - Simple
 - Sequential blocks
- Disadvantage
 - Not suitable for updates
 - Becomes ugly (loses sequentiality and balance) over time

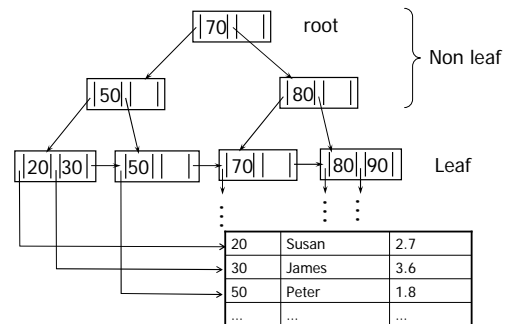
20

B+ Tree

- Most popular index structure in RDBMS
- Advantage
 - Suitable for dynamic updates
 - Balanced
 - Minimum space usage guarantee
- Disadvantage
 - Non-sequential index blocks


21

B+ Tree Example (n=3)



22


Sample Leaf Node (n=3)



- n: max # of pointers in a node
- All pointers (except the last one) point to tuples
- At least half of the pointers are used.
(more precisely, $\lceil (n+1)/2 \rceil$ pointers)

23

Sample Non-leaf Node (n=3)




- Points to the nodes one-level below
 - No direct pointers to tuples
- At least half of the ptrs used (precisely, $\lceil n/2 \rceil$)
 - except root, where at least 2 ptrs used

24

Search on B+ tree

- Find 30, 60, 70?




- Find a greater key and follow the link on the left
(Algorithm: Figure 12.10 on textbook)

25

Nodes are never too empty

- Use at least
Non-leaf: $\lceil n/2 \rceil$ pointers
Leaf: $\lceil (n+1)/2 \rceil$ pointers



26

Number of Ptrs/Keys for B+ tree

	Max Ptrs	Max keys	Min ptrs	Min keys
Non-leaf (non-root)	n	n-1	$\lceil n/2 \rceil$	$\lceil n/2 \rceil - 1$
Leaf (non-root)	n	n-1	$\lceil (n+1)/2 \rceil$	$\lceil (n-1)/2 \rceil$
Root	n	n-1	2	1

27


B+ Tree Insertion

- simple case (no overflow)
- leaf overflow
- non-leaf overflow
- new root

28

Insertion (Simple Case)


- Insert 60



29

Insertion (Leaf Overflow)

- Insert 55




- No space to store 55

30

Insertion (Leaf Overflow)

- Insert 55




- Q: After split, leaf nodes always half full?

31

Insertion (Non-leaf Overflow)

- Insert 52




Leaf overflow. Split and copy the first key of the new node

32

Insertion (Non-leaf Overflow)

- Insert 52




- Q: After split, non-leaf at least half full?

33

Insertion (New Root Node)


- Insert 25



34

Insertion (New Root Node)

- Insert 25
- Q: At least 2 ptrs at root?



35

B+ Tree Insertion

- Leaf node overflow
 - The first key of the new node is copied to the parent
- Non-leaf node overflow
 - The middle key is moved to the parent
- Detailed algorithm: Figure 12.13

36

B+Tree Deletion


- (a) Simple case (no underflow)
- (b) Leaf node, coalesce with neighbor
- (c) Leaf node, redistribute with neighbor
- (d) Non-leaf node, coalesce with neighbor
- (e) Non-leaf node, redistribute with neighbor

In the examples, $n = 4$

- Underflow for non-leaf when fewer than $\lceil n/2 \rceil = 2$ ptrs
- Underflow for leaf when fewer than $\lceil (n+1)/2 \rceil = 3$ ptrs
- Nodes are labeled as a, b, c, d, \dots

37


(a) Simple case



- Delete 25

38


(b) Coalesce with sibling (leaf)



- Delete 50

39


(b) Coalesce with sibling (leaf)



- Delete 50
 - Check underflow at a . Min 2 ptrs, currently 3

40


(c) Redistribute (leaf)



- Delete 50

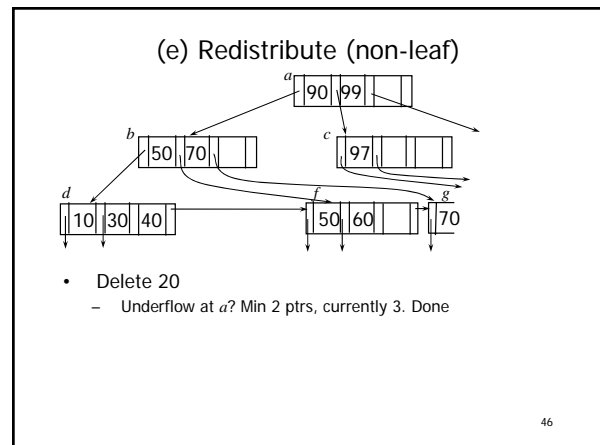
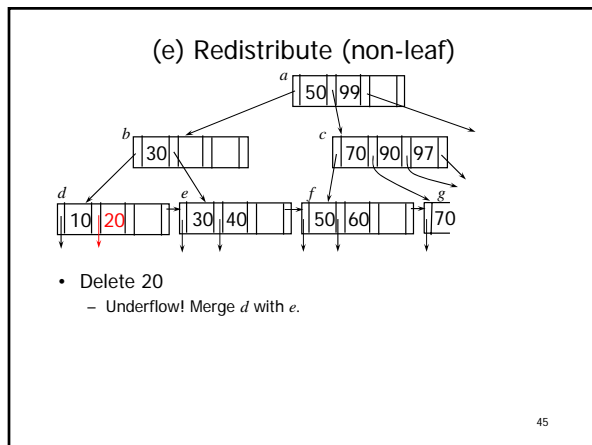

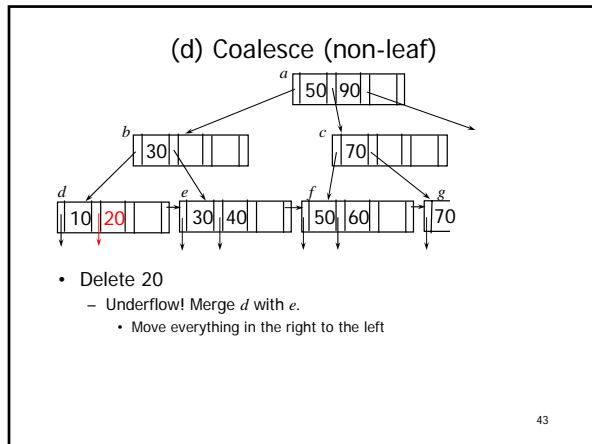
41

(c) Redistribute (leaf)



- Delete 50
 - No underflow at a . Done.

42



Important Points

- Remember:
 - For leaf node merging, we delete the mid-key from the parent
 - For non-leaf node merging/redistribution, we pull down the mid-key from their parent.
- Exact algorithm: Figure 12.17
- In practice
 - Coalescing is often not implemented
 - Too hard and not worth it

47


Where does n come from?

- n determined by
 - Size of a node
 - Size of search key
 - Size of an index pointer
- Q: 1024B node, 10B key, 8B ptr $\rightarrow n$?

48

Question on B+tree

- `SELECT *`
`FROM Student`
`WHERE sid > 60?`



49

Summary on tree index

- Issues to consider
 - Sparse vs. dense
 - Primary (clustering) vs. secondary (non-clustering)
- Indexed sequential file (ISAM)
 - Simple algorithm. Sequential blocks
 - Not suitable for dynamic environment
- B+ trees
 - Balanced, minimum space guarantee
 - Insertion, deletion algorithms

50

Index Creation in SQL

- `CREATE INDEX <indexname>`
`ON <table>(<attr>,<attr>,...)`
- Example
 - `CREATE INDEX stdix ON Student(sid)`
 - Creates a B+tree on the attributes
 - Speeds up lookup on sid

51

Primary (Clustering) Index

- MySQL:
 - Primary key becomes the clustering index
- DB2:
 - `CREATE INDEX idx ON Student(sid) CLUSTER`
 - Tuples in the table are sequenced by sid
- Oracle: Index-Organized Table (IOT)
 - `CREATE TABLE T (`
...
`) ORGANIZATION INDEX`
 - B+tree on primary key
 - Tuples are stored at the leaf nodes of B+tree
- Periodic reorganization may still be necessary to improve range scan performance

52

Next topic

- Hash index
 - Static hashing
 - Extendible hashing

53

What is a Hash Table?

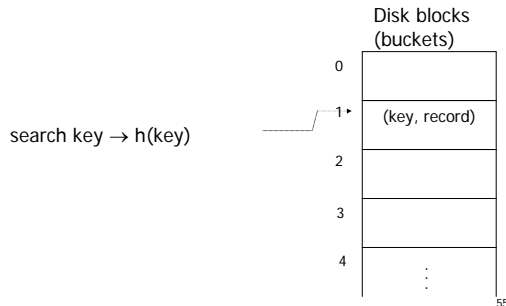
- Hash Table
 - Hash function
 - $h(k): \text{key} \rightarrow \text{integer } [0..n]$
 - e.g., $h(\text{'Susan'}) = 7$
 - Array for keys: $T[0..n]$
 - Given a key k , store it in $T[h(k)]$

$h(\text{Susan}) = 4$
 $h(\text{James}) = 3$
 $h(\text{Neil}) = 1$

0	
1	Neil
2	
3	James
4	Susan
5	


54

Hashing for DBMS (Static Hashing)



Overflow and Chaining


- Insert
 $h(a) = 1$
 $h(b) = 2$
 $h(c) = 1$
 $h(d) = 0$
 $h(e) = 1$
- Delete
 $h(b) = 2$
 $h(c) = 1$



56

Major Problem of Static Hashing


- How to cope with growth?
 - Data tends to grow in size
 - Overflow blocks unavoidable



57

Extendible Hashing (two ideas)


- (a) Use i of b bits output by hash function



58

Extendible Hashing (two ideas)

- (b) Use directory that maintains pointers to hash buckets (indirection)




59

Example

- $h(k)$ is 4 bits; 2 keys/bucket


Insert 0111



60

Example

Insert 1010 (overflow)




Increase i of the bucket. Split it.

61

Example


Insert 0000



Split bucket and increase i

62

Insert 0011 (double directory)



Split bucket, increase i ,
redistribute keys


63

Extendible Hashing: Deletion

- Two options
 - No merging of buckets
 - Merge buckets and shrink directory if possible

64

Delete 1010 (merge buckets and shrink directory)



65

Bucket Merge Condition

- Bucket merge condition
 - Bucket i 's are the same
 - First $(i-1)$ bits of the hash key are the same
- Directory shrink condition
 - All bucket i 's are smaller than the directory i


66

Questions on Extendible Hashing

- Can we provide minimum space guarantee?

67

Space Waste



68

Hash index summary

- Static hashing
 - Overflow and chaining
- Extendible hashing
 - Can handle growing files
 - No periodic reorganizations
 - Indirection
 - Up to 2 disk accesses to access a key
 - Directory doubles in size
 - Not too bad if the data is not too large

69

Hashing vs. Tree

- Can an extendible-hash index support?


```
SELECT
FROM R
WHERE R.A > 5
```
- Which one is better, B+ tree or Extendible hashing?


```
SELECT
FROM R
WHERE R.A = 5
```

70