

CS143 Notes: TRANSACTION

Book Chapters

(4th) Chapters 15, 16.1, 16.7-8, 17.1-4, 17.6

(5th) Chapters 15, 16.1, 16.7-8, 17.1-5

MOTIVATION FOR TRANSACTION

1. Crash recovery

- \langle eg, Transfer \$1M from Susan to Jane \rangle (example slide)
 - S_1 : UPDATE Account SET balance = balance - 1000000 WHERE owner = ‘Susan’
 - S_2 : Update Account SET balance = balance + 1000000 WHERE owner = ‘Jane’
 - System crashes after S_1 but before S_2 . What now?

2. Concurrency


- We do not want to allow oncurrent access from multiple clients. We do not want to “lock out” the DBMS until one client finishes
 \langle explain with client/server diagram \rangle
- Can allow parallel execution while avoiding any potential problems from concurrency?
(we will see concurrency problem examples soon).

TRANSACTION AND “ACID” PROPERTY

- TRANSACTION: A sequence of SQL statements that are executed as a “unit”
- ACID PROPERTY OF TRANSACTION: Atomicity, Consistency, Isolation, Durability

1. Durability

- If a transaction committed, all its changes remain permanently even after system crash
- COMMENTS: Not very easy because some changes may be reflected only in memory for performance reasons
- 2. Atomicity: “ALL-OR-NOTHING”
 - Either ALL OR NONE of the operations in a transaction is executed.
 - If the system crashes in the middle of a transaction, all changes by the transaction are “undone” during recovery.
- 3. Consistency: If the database is in a consistent state before a transaction, the database is in a consistent state after the transaction
- 4. Isolation: Even if multiple transactions are executed concurrently, the result is the same as executing them in some sequential order.
 - Each transaction is unaware of (is isolated from) other transaction running concurrently in the system
(explain by time line diagram)



- DBMS guarantees the ACID property for all transactions
 - With minor caveats that will be discussed later.
- **Q:** How can the database system guarantee these? Any ideas?

DECLARING A TRANSACTION IN SQL

- Two important commands:
 - COMMIT: All changes made by the transaction is stored permanently
 - ROLLBACK: Undo all changes made by the transaction
- AUTOCOMMIT MODE
 1. With AUTOCOMMIT mode OFF
 - Transaction implicitly begins when any data in DB is read or written
 - All subsequent read/write is considered to be part of the same transaction
 - A transaction finishes when COMMIT or ROLLBACK statement is executed
(explain using time line diagram)



2. With AUTOCOMMIT mode ON

- Every SQL statement becomes one transaction

- Setting Autocommit mode:

- In DB2: UPDATE COMMAND OPTIONS USING c ON/OFF (default is on)
- In Oracle: SET AUTOCOMMIT ON/OFF (default is off)
- In MySQL: SET AUTOCOMMIT = {0|1} (default is on. InnoDB only)
- In JDBC: connection.setAutoCommit(true/false) (default is on)
- In Oracle and MySQL, BEGIN temporarily disables autocommit mode until COMMIT or ROLLBACK

TWO QUESTIONS ON TRANSACTION AND CONCURRENCY

1. What execution orders are “good”?

- We first need to understand what execution orders are okay
 - Serializability theory

2. How can we enforce “good” execution order?

- Concurrency control mechanism

Topics of next discussion

“GOOD” SCHEDULE

- $\langle \text{ex, salary increase} \rangle$
(Show example slides)
 - T_1 :
 - * UPDATE Employee SET salary = salary + 100 WHERE name = ‘Susan’
 - * UPDATE Employee SET salary = salary + 100 WHERE name = ‘Jane’
 - T_2 :
 - * UPDATE Employee SET salary = salary * 2 WHERE name = ‘Susan’
 - * UPDATE Employee SET salary = salary * 2 WHERE name = ‘Jane’
- Constraint: Susan’s salary = Jane’s salary

Internally, DBMS performs the following operations

A: Susan.salary, B: Jane.salary

- T_1 : Read(A), A=A+100, Write(A), Read(B), B=B+100, Write(B)
- T_2 : Read(A), A=A*2, Write(A), Read(B), B=B*2, Write(B)
- SCHEDULE: The chronological order that instructions are executed in the system
(Show example schedules and show that some schedules are okay and some are not)

– Schedule A:

T_1	T_2
Read(A); A=A+100	
Write(A)	
Read(B); B=B+100	
Write(B)	
	Read(A); A=A*2
	Write(A)
	Read(B); B=B*2
	Write(B)

– Schedule B:

T_1	T_2
	Read(A); A=A*2
	Write(A)
	Read(B); B=B*2
	Write(B)
Read(A); A=A+100	
Write(A)	
Read(B); B=B+100	
Write(B)	

- Schedule C:

T_1	T_2
Read(A); A=A+100 Write(A)	Read(A); A=A*2 Write(A)
Read(B); B=B+100 Write(B)	Read(B); B=B*2 Write(B)

- * **Q:** We get the same or “equivalent” result for schedule A and C. Is it just a coincidence? Is there some reason behind it? (More discussion later...)

- Schedule D:

T_1	T_2
Read(A); A=A+100 Write(A)	Read(A); A=A*2 Write(A)
Read(B); B=B+100 Write(B)	Read(B); B=B*2 Write(B)

- * **Q:** Both C and D executes T_1 and T_2 concurrently. Why is one okay but not the other? Any fundamental reason?

- Claim: Two schedules do the “same” thing if

1. all actions in the two transactions read the same value
2. they write the same final results to the database

- Intuitively, as long as the transactions read the same values, they will take the same actions. As long as READs and WRITEs are the same transactions do the same thing.

⟨a simple example⟩

T_1	T_2
$V = r(A)$	$W = r(A)$
$V = V + 100$	$W = W + 200$
$w(A, V)$	$w(A, W)$

- * T_1 and T_2 take actions based on input values
 - As long as T_1 and T_2 read the same value, T_1 and T_2 do the same thing
- * We also care about the final output to the database, so the writing to the database is important

- In writing schedules, we just focus on read/write operations

- * What is really important is “read” and “write” actions. All other things depend on these two actions.
- * As long as T_1 and T_2 read and write the same values, two schedules are “equivalent”

- Notation for schedule

$$\begin{array}{c|c}
 - & \begin{array}{c} T_1 \\ \hline V = r(A) \\ V = V + 100 \\ \\ \\ w(A,V) \end{array} & \begin{array}{c} T_2 \\ \hline \\ \\ W = r(A) \\ W = W + 200 \\ w(A,W) \end{array} \\
 \hline
 \end{array}$$

S : $r_1(A) \ r_2(A) \ w_2(A) \ w_1(A)$

* subscript 1 means transaction 1

* $r(A)$ means read A

* $w(A)$ means write to A

SERIAL SCHEDULE

- Consider schedule A: $S_a = \frac{r_1(A) \ w_1(A) \ r_1(B) \ w_1(B)}{T_1} \ \frac{r_2(A) \ w_2(A) \ r_2(B) \ w_2(B)}{T_2}$

– **Q:** Is this schedule “good”?

- SERIAL SCHEDULE: all operations in any transaction are performed without any interleaving
 - Definitely a good schedule. The result is definitely good as long as individual transactions are correct.

CONFLICT SERIALIZABLE SCHEDULE

- We got the same result for S_a and S_c . Is there a fundamental reason behind this?

$$S_a = \frac{r_1(A) \ w_1(A) \ r_1(B) \ w_1(B)}{T_1} \ \frac{r_2(A) \ w_2(A) \ r_2(B) \ w_2(B)}{T_2}$$

$$S_c = \frac{r_1(A) \ w_1(A)}{T_1} \ \frac{r_2(A) \ w_2(A)}{T_2} \ \frac{r_1(B) \ w_1(B)}{T_1} \ \frac{r_2(B) \ w_2(B)}{T_2}$$

– **Q:** Is S_c a serial schedule?

– **Q:** When we swap $w_2(A)$ and $r_1(B)$, do we get the “same” result?

$$S_c: S_c = r_1(A) \ w_1(A) \ r_2(A) \ w_2(A) \ r_1(B) \ w_1(B) \ r_2(B) \ w_2(B)$$

$$S_c': S_c' = r_1(A) \ w_1(A) \ r_2(A) \ \underline{r_1(B) \ w_2(A)} \ w_1(B) \ r_2(B) \ w_2(B)$$

- CONFLICTING ACTIONS vs NON-CONFLICTING ACTIONS:

- Non-conflicting actions: A pair of actions that do not change the result when we swap them
- \langle examples \rangle
 - * **Q:** $r_1(A) \ w_2(B)$? conflicting?
 - * **Q:** $r_1(A) \ r_2(A)$? conflicting?
 - * **Q:** $w_1(A) \ r_2(A)$? conflicting?
 - * **Q:** $w_1(A) \ w_2(A)$? conflicting?
- Two actions are CONFLICTING if (1) they involve the same objects/variables and (2) at least one of them is WRITE.
- Swapping non-conflicting actions do not change results
- **Q:** Can we swap only non-conflicting actions of S_c to transform it into S_a ?
 - * COMMENTS: S_a and S_c are guaranteed to produce the same results if we can.
- \langle swap actions in S_c one by one to transform it \rangle

- CONFLICT EQUIVALENCE

- DEFINITION: S_1 is CONFLICT EQUIVALENT to S_2 if S_1 can be rearranged into S_2 by a series of swaps of non-conflicting actions.
- When S_1 and S_2 are conflict equivalent, they read and write same values. Thus, the results are always the same.

- CONFLICT SERIALIZABILITY

- Intuition:
 - * S_a and S_c always generate the same results (because they are conflict equivalent)
 - * S_a is guaranteed to generate “good” result because it is SERIAL schedule
 - * Thus, S_c is guaranteed to generate “good” result.
- FORMAL DEFINITION: S_1 is CONFLICT SERIALIZABLE if it is conflict equivalent to some serial schedule.
- A conflict serializable schedule may not be a serial schedule, but it is still a “good” schedule
- Example:
 - * **Q:**

$$S_1 = r_1(A) \ w_2(A) \ r_1(B)$$

$$S_2 = r_1(A) \ r_1(B) \ w_2(A) \text{ conflict equivalent?}$$

- **Q:** What does it mean?
- * **Q:** S_1 conflict serializable?
- **Q:** What does it mean?
- Consider schedule D that generates a “wrong” result
 Schedule D: $S_c = \frac{r_1(A) \ w_1(A)}{T_1} \ \frac{r_2(A) \ w_2(A)}{T_2} \ \frac{r_2(B) \ w_2(B)}{T_2} \ \frac{r_1(B) \ w_1(B)}{T_1}$
 - **Q:** How should we reorder actions to make it serial?
 - * Move the second block of T_1 after the first block of T_1
 - * Move the first block of T_1 before the second block of T_1
 - **Q:** Can we move $r_1(B)$ before $r_2(A)$ and get the same result?
 - **Q:** Can we move $r_1(A)$ before $r_1(B)$ and get the same result?
 - **Q:** Is Schedule D conflict serializable?
 - **Q:** In general, how can we know whether a schedule is conflict serializable?
 - Intuition: Conflicting actions cause precedence relationship between transactions:
 - * For example, in schedule D,
 - $w_2(B), \dots, r_1(B)$: $T_2 \rightarrow T_1$ because T_1 reads the value from T_2
 - eg. $r_1(A), \dots, w_2(A)$: $T_1 \rightarrow T_2$ because T_1 reads the value before T_2
 - **Q:** What does $T_1 \rightarrow T_2$ and $T_2 \rightarrow T_1$ in S_d mean?
 - If there is a cycle in the precedence relationship, there is no equivalent serial schedule.

- PRECEDENCE GRAPH $P(S)$

- Nodes: transactions in S
- Edges: $T_i \rightarrow T_j$ if
 1. $p_i(A), q_j(A)$ are actions in S
 2. $p_i(A)$ precedes $q_j(A)$
 3. At least one of p_i, q_j is a write
- THEOREM: $P(S)$ is acyclic $\Leftrightarrow S$ is conflict serializable
- $\langle \text{eg, } P(S) \text{ for } S = w_3(A) w_2(C) r_1(A) w_1(B) r_1(C) w_2(A) r_4(A) w_4(D) \rangle$
 - * **Q:** Is S conflict serializable?

- **Q:** does $P(S_1) = P(S_2)$ imply S_1 and S_2 are conflict equivalent?

* $\langle \text{eg, } S_1 = w_1(A) r_2(A) w_2(B) r_1(B) \text{ and } S_2 = r_2(A) w_1(A) r_1(B) w_2(B) \rangle$

- SUMMARY SO FAR:

- good schedule: conflict serializable schedule
- conflict serializable \Leftrightarrow acyclic precedence graph
- Now more issues related to the meaning of COMMIT

RECOVERABLE SCHEDULE

- $\langle \text{example for recoverable schedule} \rangle$

$T_1: r(A) w(A)$

$T_2: r(A) w(A)$

T_1	T_2
$r_1(A)$ $w_1(A)$	$r_2(A)$ $w_2(A)$ Can we COMMIT here?
What if ROLLBACK?	

- **Q:** Committed transaction may to be aborted! How can we avoid this scenario?

- **RECOVERABLE SCHEDULE**

- Schedule S is **RECOVERABLE** if T_j reads a data item written by T_i , the COMMIT operation of T_i appears before the COMMIT operation of T_j
- Notation:
 - * c_i : COMMIT by transaction T_i
 - * a_i : ROLLBACK by transaction T_i (abort)
- **Q:** recoverable schedule?

T_1	T_2
$r_1(A)$ $w_1(A)$	$r_2(A)$ $w_2(A)$ c_2
c_1	

- **Q:** recoverable schedule?

T_1 :	T_2
$r_1(A)$ $w_1(A)$	$r_2(A)$ $w_2(A)$ c_2
c_1	

CASCADELESS SCHEDULE

- **Q:** What happens to T_2 ?

T_1	T_2
$r_1(A)$	
$w_1(A)$	
	$r_2(A)$
	$w_2(A)$
a_1	

- CASCADING ROLLBACK: A single transaction abort leads to a series of transaction rollback

- **Q:** What is the main cause of cascading rollback?

- DIRTY READ: reading an output of an uncommitted transaction

- **Q:** How can we avoid cascading rollback?

- CASCADELESS SCHEDULE

- Schedule S is CASCADELESS if T_j reads a data item written by T_i , the COMMIT operation of T_i appears before the READ operation of T_j
- Note the difference between RECOVERABLE and CASCADELESS schedule

RELATIONSHIP BETWEEN SCHEDULES

- **Q:** What is the relationship between recoverable and cascadeless schedule?
(Vann diagram)
- **Q:** What is the relationship between serial and cascadeless schedule?
(Vann diagram)
- **Q:** What is the relationship between serial and conflict-serializable schedule?
- **Q:** What is the relationship between conflict-serializable and recoverable schedule?

- **Q:** If a schedule is conflict serializable, is it recoverable?

⟨example⟩

T_1	T_2
$w_1(A)$	
	$r_2(A)$
	c_2
c_1	

- **Q:** If a schedule is recoverable, is it conflict serializable?

⟨example⟩

T_1	T_2
	$r_2(A)$
$w_1(A)$	
	$r_2(A)$
c_1	
	c_2

- **Example S:** $w_1(A) \ w_2(A) \ w_1(B) \ r_2(B) \ c_1 \ c_2$

- **Q:** Is it serial?
- **Q:** Is it conflict-serializable?
- **Q:** Is it recoverable?
- **Q:** Is it cascadeless?
- **Q:** How can we make it cascadeless?

WHAT TO REMEMBER

- Conflict serializable schedule
- Recoverable schedule
- Cascadeless schedule
- We want either
 - Conflict serializable + Recoverable
 - Conflict serializable + Cascadeless

LOCKING PROTOCOL

- Main question: How can we achieve serializable and cascadeless schedule?
- Let us think how we may achieve serializable and cascadeless schedule.

T_1	T_2
$r_1(A)$	
$w_1(A)$	
	$r_2(A)$
	$w_2(A)$
a_1	

- **Q:** Why do we have cascading rollback? How can we avoid it?
- **Q:** How can we ensure T_2 does not read T_1 's write until T_1 commits?

Rigorous Two Phase Locking Protocol (R2PL)

- Basic idea: Avoid dirty read by “locking” modified values until commit.
 1. Before T_1 writes, T_1 obtains a lock on A.
 2. T_1 releases the lock only when T_1 commits.
 3. When T_1 holds the lock on A, T_2 cannot access A.
- Three rules for Rigorous Two Phase Locking Protocol
 - Rule (1): T_i has to lock tuple t_k before any read/write
 - Rule (2): When T_i is holding the lock on t_k , T_j cannot obtain the lock on t_k (for $j \neq i$)
 - Rule (3): Release all locks at commit
(explain using lock accumulation diagram)
- IMPORTANT THEOREM
 - Rigorous 2PL ensures a conflict-serializable and cascadeless schedule.
- RIGOROUS 2PL SCHEDULE
 - A schedule that can be produced by rigorous 2PL protocol
 - **Q:** Is there any conflict-serializable and cascadeless schedule that cannot be produced by R2PL?

Two Phase Locking Protocol (2PL)

- Less strict locking protocol than rigorous 2PL
 - Rule (1): T_i lock a tuple before any read/write
 - Rule (2): If T_i holds the lock on A, T_j cannot access A ($j \neq i$)
 - Rule (3): Two stages:
 - * (a) growing stage: T_i may obtain locks, but may not release any lock
 - * (b) shrinking stage: T_i may release locks, but may not obtain any lock

⟨explain using lock accumulation diagram⟩
- IMPORTANT THEOREM
 - 2PL ensures a conflict-serializable schedule
- 2PL SCHEDULE
 - A schedule that can be produced by 2PL protocol.
 - **Q:** Is there any conflict-serializable schedule that cannot be produced by 2PL?
- **Q:** What's the relationship of R2PL and serializable/cascadeless schedule?
⟨explain using Vann Diagram⟩
- **Q:** What's the relationship of 2PL and serializable/cascadeless schedule?
⟨explain using Vann Diagram⟩

RECOVERY AND LOGGING

- Motivation for logging. Consider $T : r(A)w(A)r(B)w(B)$.
 - **Example 1:** $S = r(A)w(A)r(B)a$. What should we do? How do we get the old value of A ?
 - **Example 2:** $S = r(A)w(A) !!!CRASH!!!$ What should DBMS do when it reboots?
 - **Example 3:** $S = r(A)w(A)r(B)w(B)c$. New A and B values are “cached” in main memory for performance reasons. Can DBMS commit T without writing the new values permanently to the disk?
⟨main-memory and disk diagram⟩
- Rules for log-based recovery
 1. For every action DBMS performs, a “log record” for the action should be generated.
 - $\langle T_i, \text{start} \rangle$
 - $\langle T_i, X_j, \text{old-value}, \text{new-value} \rangle$
 - $\langle T_i, \text{commit} \rangle$
 - $\langle T_i, \text{abort} \rangle$
 2. Log record should be written to disk BEFORE the actual data is written to the disk.
 - $\langle T_i, A, 5, 10 \rangle$ should be written before the new A value 10 is written to the data block.
 3. Before commit of T_i , all log records for T_i are written to disk (including commit).
 - The actual data block may or may not be written to disk at commit.
 4. During abort, DBMS gets old values from the log
 5. During recovery, DBMS “re-executes” all actions in the committed transactions and “rolls back” all actions in the non-committed transactions.

- **Example:**

⟨Explain log records line by line⟩

$A: 100, B: 100, C: 100$

T_1	T_2	Log
Read(A); $A=A-50$		1 $\langle T_1, \text{start} \rangle$
Write(A)		2 $\langle T_1, A, 100, 50 \rangle$
	Read(C); $C=C*2$	3 $\langle T_2, \text{start} \rangle$
	Write(C)	4 $\langle T_2, C, 100, 200 \rangle$
	Commit	5 $\langle T_2, \text{commit} \rangle$
Read(B); $B=B+50$		
Write(B)		6 $\langle T_1, B, 100, 150 \rangle$
Commit		7 $\langle T_1, \text{commit} \rangle$

- **Q:** What should DBMS do during recovery when it sees up to log record 4?
- **Q:** What should DBMS do during recovery when it sees up to log record 5?
- **Q:** What should DBMS do during recovery when it sees up to log record 7?

SQL ISOLATION LEVELS

- Motivation: In some cases, we may not need full ACID. We may want to allow some “bad” schedule to achieve more concurrency
 - SQL isolation levels allow a few “bad” scenarios for more concurrency
 - * dirty read, non-repeatable read, phantom
 - We go over three scenarios in which “relaxing” the strict ACID may be desirable for some applications

- ⟨explain the isolation levels through examples and fill in the table⟩

isolation level	dirty read	nonrepeatable read	phantom
read uncommitted			
read committed			
repeatable read			
serializable			

- DIRTY READ may be OK

- ⟨example⟩
 - * T_1 : UPDATE Employee SET salary = salary + 100
 - * T_2 : SELECT salary FROM Employee WHERE name = ‘John’
- **Q:** Under ACID, once T_1 update John’s salary, can T_2 read John’s salary?

* Sometimes, it may be okay for T_2 to proceed.

- DIRTY READ: a transaction reads uncommitted values
- “READ UNCOMMITTED” isolation level allows dirty read.
(Fill in the dirty read column)

- NON-REPEATABLE READ may be OK

- ⟨example⟩
 - * T_1 : UPDATE Employee SET salary = salary + 100 WHERE name = ‘John’
 - * T_2 : (S_1) SELECT salary FROM Employee WHERE name = ‘John’
 - ...
 - (S_2) SELECT salary FROM Employee WHERE name = ‘John’
- **Q:** Under ACID, can we get different values for S_1 and S_2 ?

* Sometimes it may be okay to get different values

- NON-REPEATABLE READ: When T_i reads the same row multiple times, T_i may get different values
- “READ UNCOMMITTED” or “READ COMMITTED” isolation levels allow NON-REPEATABLE READ.
(Fill in the non-repeatable read column)

- PHANTOM may be OK

- \langle example \rangle
 - * Initially, $\text{SUM}(\text{Employee.salary}) = \$100,000$
 - * T_1 : `INSERT INTO Employee (e1, 1000), (e2, 1000)`
 - * T_2 : `SELECT SUM(salary) FROM Employee`
- **Q:** Under ACID, what may T_2 return?

- * Sometimes, it may be OK for T_2 to return \$101,000

- **Q:** Under REPEATABLE READ, what if T_2 is

```
SELECT SUM(salary) FROM Employee
...
SELECT SUM(salary) FROM Employee
```

What can T_2 return?

- PHANTOM: When new tuples are inserted, once some of them are seen by statements, or only some statements see the newly inserted tuples.
- Except for “SERIALIZABLE” isolation level, PHANTOM is always allowed.

- MIXED ISOLATION LEVELS

- \langle example on mixed isolation levels \rangle
 - * T_1 : `UPDATE Employee SET salary = salary + 100`
`ROLLBACK`
 - * T_2 : `SELECT salary FROM Employee WHERE name = 'John'`
- **Q:** T_1 - SERIALIZABLE, T_2 - SERIALIZABLE. What may T_2 return?

- **Q:** T_1 - SERIALIZABLE, T_2 - READ UNCOMMITTED. What may T_2 return?

- COMMENTS:

- * Only when all transactions are serializable, we guarantee ACID.
 - * The isolation level is in the eye of the beholding transaction.

- READ ONLY TRANSACTION

- Many, many transactions are read only.
- By declaring a transaction as READ ONLY, we can help DBMS to optimize for more concurrency
- SQL ISOLATION LEVEL DECLARATION
 - SET TRANSACTION options
 - access mode: READ ONLY / READ WRITE (default: READ WRITE)
 - isolation level: ISOLATION LEVEL
 - * READ UNCOMMITTED
 - * READ COMMITTED (Oracle default)
 - * REPEATABLE READ (MySQL, DB2 default)
 - * SERIALIZABLE
 - e.g) SET TRANSACTION READ ONLY, REPEATABLE READ
 - * READ UNCOMMITTED cannot be READ WRITE
 - * Needs to be declared before EVERY transaction for non-default settings

OPTIONAL MATERIALS

Proof of P(S) is acyclic \Leftrightarrow S is conflict serializable

- Lemma: S_1, S_2 conflict equivalent $\Rightarrow P(S_1) = P(S_2)$
- Proof:
 - Assume $P(S_1) \neq P(S_2)$
 - \Rightarrow There exists $T_i \rightarrow T_j$ in S_1 but not in S_2
 - $\Rightarrow S_1 = \dots p_i(A) \dots q_j(A) \dots$ ($p_i(A), q_j(A)$: conflicting actions)
 - $S_2 = \dots p_j(A) \dots p_i(A) \dots$
 - $\Rightarrow S_1$ and S_2 are not conflict equivalent!
- Proof 1: $P(S)$ is acyclic \Leftarrow S is conflict serializable
 - Assume S is conflict serializable
 - \Rightarrow There exists S_s where S_s and S are conflict equivalent
 - $\Rightarrow P(S_s) = P(S)$ from Lemma
 - $\Rightarrow P(S)$ acyclic because $P(S_s)$ is acyclic
- Proof 2: $P(S)$ is acyclic \Rightarrow S is conflict serializable
 - Assume $P(S)$ is acyclic. Transform S as follows:
 1. Take T_1 to be a transaction with no incident edges
 - **Q:** Is there always a transaction with no incident edge?
 2. Move all T_1 actions to the front

$$S_1 = \dots q_j(A) \dots \quad p_1(A) \dots$$

\swarrow
 3. We now have $S = \langle T_1 \text{ actions} \rangle \langle \dots \text{rest} \dots \rangle$
 4. Repeat the above steps to serialize the rest

SHARED & EXCLUSIVE LOCK

- $\langle \text{example} \rangle$

T_1	T_2
$r(A)$	
$r(A)$	$r(A)$
	$r(B)$

- **Q:** Is it possible for our R2PL protocol?
- **Q:** Is it conflict serializable?

- COMMENTS: $r(A)$ and $r(A)$ do not conflict. Only $w(A)$ causes conflict. We should have granted locks to T_1 and T_2 for “reading”

- SHARED & EXCLUSIVE LOCK

- SHARED LOCK:
 - * lock for read
 - * multiple transactions can obtain the same shared lock
- EXCLUSIVE LOCK
 - * lock for write
 - * If T_i holds an exclusive lock for A, no other transaction can obtain a shared/exclusive lock on A.
- Separate locks for read and write
 - * Before read on A, T_i requests a SHARED lock on A
 - * Before write on A, T_i requests an exclusive lock on A
 - * Everything else is the same as before.

- COMPATIBILITY MATRIX

	shared	exclusive
shared		
exclusive		

- Rigorous 2PL with shared lock \rightarrow conflict serializable and cascadeless
2PL with shared lock \rightarrow conflict serializable

PHANTOM PROBLEM

- \langle Phantom example schedule slide \rangle

eid	salary	// table is sequenced by eid
e3	1000	

- T_1 : SELECT SUM(salary) FROM Employee
- T_2 : INSERT INTO Employee VALUES (e1, 500), (e5, 500)
 - * e1 is inserted before e3
 - * e5 is inserted after e5
 - * T_1 maintains a cursor to scan the table
- **Q:** What result does T_1 return?
 - * Follow the schedule sequentially
- **Q:** Does the schedule follow rigorous R2PL?

- **Q:** Does the schedule provide ACID?

- **Q:** Why do we get this result?
 - * T_1 scans the “entire table” not just e3.
 - * When T_1 is at e3, either
 1. T_1 should have read e1, or
 2. T_2 should not be allowed to insert e1.
 → T_1 has to worry about “non-existing” e1 tuple: PHANTOM PHENOMENON
- **Q:** How can we avoid this?

- INSERT LOCK on table
 1. Before insertion, T_i gets an INSERT LOCK for the table
 2. Before scanning a table, T_i gets a INSERT LOCK for the table
 3. NOTE: T_i should still obtain a lock on each tuple that it it reads/writes on
 - INSERT LOCK simply prevents the insertion of a new tuple into a TABLE. Individual tuples should be “protected” by their own locks before read/write.

DEADLOCK

- DEADLOCK: multiple transactions wait for each other without making any progress
 - 2PL may cause deadlock in certain cases.

⟨example⟩

T_1 : $r(A)w(B)$

T_2 : $r(B)w(A)$

T_1	T_2
l(A)	
r(A)	
	l(B)
	r(B)
l(B)?	
w(B)	
	l(A)?
	w(A)

- **Q:** Can T_1 and T_2 progress?
 (explain using wait-graph)

- **Q:** What should we do in this case to get out of the deadlock?

- COMMENTS:

1. Most DBMS runs deadlock detection algorithm.
2. If there is a deadlock, roll-back transactions until the deadlock is broken

SQL ISOLATION LEVEL Implementation mechanisms

- Tuple lock:
 - S: shared lock on tuple t before read
 - X: exclusive lock on tuple t before write
- Table INSERT lock:
 - S: shared insert lock on table T before read/write
 - X: exclusive insert lock on table T before insert

(explain implementaion mechanism filling in the following table)

	read		insert		write	
isolation level	t	I	t	I	t	I
serializable						
repeatable read						
read committed						
read uncommitted						

- first discuss serializable:
 - * We learned the mechanisam already. Rigorous 2PL with insert lock.
- Then discuss first on read for every other isolation level
 - * start with read uncommitted, then read committed
 - * **Q:** For “read uncommitted” transaction,
 - READ LOCK before read?
 - INSERT LOCK before read?

- then insert
 - and then write
- COMMENTS:
 1. READ COMMITTED may release the shared lock immediately after read
 2. REPEATABLE READ holds the shared lock until it does not access the tuple any more
 3. All exclusive locks are held until COMMIT