

CS143 Notes: Views & Authorization

Views

- What is a view?
 - A “virtual” table created on top of other “real” tables
 - Almost the same as a “real” table except that
 - * the tuples are computed on the fly using “real” tables.
 - * a view does not really “exist”.
- Syntax and example:
 - `CREATE VIEW ViewName(A1, A2, ...) AS Query`
 - * Attribute lists are optional
 - * Example: SidNameAddr view with (sid, name, address) from Student
 - Views can be used in a query like:

```
SELECT *
FROM SidNameAddr S, Enroll E
WHERE S.sid = E.sid
```

 - * The system automatically rewrites the query using the SidNameAddr view definition
 - Views can be created on top of other views

```
CREATE VIEW NameAddr AS
  SELECT name, addr
  FROM SidNameAddr
```
 - **Q:** MultiClass: View of students (sid, name) who take more than one class?
 - **Q:** Why use views?

- Three-level vision of database:
 - * Virtual database: Views $V := \text{ViewQuery}(R_1, R_2, \dots, R_n)$
built on top of ...
 - * Conceptual database: Tables (Relations)
built on top of ...
 - * Physical database: Pages on disk

Modifying Views

- Updates on views are allowed (under certain conditions)
- **Q:** How can we “update” a view when it does not exist?
 - **Q:** UPDATE SidNameAddr SET Name = ‘James’ WHERE sid = 301?
 - * Modification to a view is “translated” into a modification to the underlying table
 - **Q:** INSERT INTO SidNameAddr VALUES (305, ‘Peter’, ‘1234 Westwood’)?
 - * Missing columns are filled with the DEFAULT value (or NULL)
 - **Q:** INSERT INTO SameAddr VALUES (‘Tony’, ‘Joshua’)?


```
CREATE VIEW SameAddr AS
  SELECT S1.name, S2.name
  FROM Student S1, Student S2
  WHERE S1.addr = S2.addr AND S1.sid > S2.sid
```
 - **Q:** UPDATE AvgGPA SET a = 3.0?


```
CREATE VIEW AvgGPA(a) AS SELECT avg(GPA) FROM Student;
```
- For some views, update may not make any sense
 - * Precise conditions for updatable views are very complicated
 - may involve keys, equality conditions, etc.
- SQL2 uses very conservative conditions. View must be defined as:
 - * SELECT on a single table T, without DISTINCT
 - * Subqueries in WHERE must not refer to T

- * Attributes of T not projected in view allowed to be NULL or default
- * No aggregation

- **Q:** INSERT INTO Student17 VALUES (403, 'Peter', '123 Olympic', 3.0, 20) ?

```
CREATE VIEW Student17 AS
SELECT *
FROM Student
WHERE age = 17
```

- **Q:** Is the new tuple in Student17?

- WITH CHECK OPTION

- * syntax: CREATE VIEW ... AS ... WITH CHECK OPTION
- * check INSERT/UPDATE to ensure the new tuple is still in the view
- * reject the statement if not

- **Q:** What will happen if we drop HonorStudent view?

```
YoungHonorStudent
  ↑
HonorStudent
  ↑
Student
```

- DROP ... [CASCADE | RESTRICT]

- * CASCADE (default): drop anything that references the view
- * RESTRICT: drop statement fails if the view is referenced by other views or integrity constraints

Materialized Views

- Some DBMS allows to “precompute” or “materialize” a view
- Example: MultiClass view again. Students who take multiple classes

```
CREATE VIEW MultiClass AS
SELECT sid, name
FROM Student, Enroll
WHERE Student.sid = Enroll.sid
GROUP BY Student.sid
HAVING COUNT(*) > 1
```

- **Q:** Why do we want to we materialize this view?

– **Q:** Why don't we always materialize views?

* **Q:** When should we refresh MultiClass?

- refresh of materialized view can be costly
- incremental refresh of materialized view is sometimes difficult

– **Q:** What views to materialize? Both? Just one? Pros and cons of each?

```
CREATE VIEW MultiClass AS
  SELECT sid, name
  FROM Student, Enroll
  WHERE Student.sid = Enroll.sid
  GROUP BY Student.sid
  HAVING COUNT(*) > 1
```

```
CREATE VIEW StudDeptCount AS
  SELECT sid, name, dept, COUNT(*)
  FROM Enroll
  GROUP BY sid, dept
```

* Deciding views to materialize is a difficult optimization problem

- Many commercial DBMS supports materialized views
 - used for “data warehouse” for OLAP (online analytical processing) queries
 - limited support for incremental refresh
 - materialized view selection is a difficult optimization problem

Authorization

- Make sure users only see what they are allowed to see
- Do not let an unauthorized user to modify database

Privileges

- For a relation R and user U, U may be authorized for:
 - SELECT ON R
 - INSERT(A₁, A₂, ..., A_n) ON R
 - * the rest of the columns should take NULL or DEFAULT
 - UPDATE(A₁, A₂, ..., A_n) ON R
 - DELETE ON R
- GRANT <privileges> ON <R> TO <users> [WITH GRANT OPTION]
 - <privileges>: SELECT, INSERT, ... separated by commas (or ALL PRIVILEGES)
 - <R>: table, view, ...
 - <users>: list of users/groups, or PUBLIC
 - more about WITH GRANT OPTION later

EXAMPLE: Grant SELECT privilege on Student to cs143

- **Q:** When will it be useful to limit insert privileges to certain columns?

- EXAMPLES

- UPDATE Student
SET GPA = 4.0
WHERE sid IN (SELECT sid FROM Enroll WHERE dept = 'CS')
 - * **Q:** What privileges are needed for this statement?
- SELETE FROM Student
WHERE sid NOT IN (SELECT sid FROM Enroll)
 - * **Q:** What privileges are needed for this statement?

Managing privileges

- **Q:** Who can grant privileges?
 - database administrator (DBA in oracle, SUPER in MySQL)
 - owner (= creator) of the table/view/...
- GRANT ... TO $\langle u_1 \rangle$ WITH GRANT OPTION
 - User u_1 can grant the same or less privilege to others
- Authorization graph.
 - Nodes: users
 - Edges:
 - * When u_i grants privilege to u_j , an edge is added from u_i to u_j .
 - * When u_1 creates a table/view ..., edge is created from DBA to u_1
 - Example: What does the following authorization graph mean?

$\text{DBA} \rightarrow u_1 \rightarrow u_2 \rightarrow u_3$

- Revoking privileges
 - REVOKE $\langle \text{privileges} \rangle$ ON $\langle R \rangle$ FROM $\langle \text{users} \rangle$ [CASCADE | RESTRICT]
 - * EXAMPLE: revoke SELECT privilege on Student from u_2
 - CASCADE/REVOKE option in REVOKE statement
 - * **Q:** Given the following authorization graph, what should happen when u_1 revokes privilege from u_2 ?


$\text{DBA} \rightarrow u_1 \rightarrow u_2 \rightarrow u_3$

- CASCADE (default): revoke all privileges passed on by the user
- RESTRICT: reject the REVOKE command

- * **Q:** What if u_3 got privilege from u_4 as well?

$\begin{array}{ccccccc} \text{DBA} & \rightarrow & u_1 & \rightarrow & u_2 & \rightarrow & u_3 \\ & & & & \searrow & & \nearrow \\ & & & & u_4 & & \end{array}$

- **Q:**



If DBA revokes privilege from u_2 , what edges should we remove? $u_1 \rightarrow u_2$? $u_3 \rightarrow u_2$?

NOTES:

- SQL allows u_1 to revoke privilege from u_2 only if u_1 granted the privilege to u_2 . This restriction is to avoid privilege revocation ambiguity.
- Unfortunately this restriction is too strict in practice, so most commercial systems do not enforce this. In case of ambiguity, they just do whatever is “reasonable” according to their policy.
- **Q:** What privileges for
`CREATE TABLE test(a int, b int),`
`CREATE VIEW ...?`
 - Unfortunately, no standard in SQL
 - * MySQL: CREATE, DROP, ALTER
 - * Oracle: CREATE TABLE, CREATE VIEW, DROP ANY TABLE, ...
 - * DB2: CREATETAB (table creation), CREATEALIAS (view creation), DROP, ALTER, ...
- Other privileges in later SQL standards
 - SQL92: REFERENCES (reference in constraints), USAGE (domain)
 - SQL99: TRIGGER, EXECUTE (function call), and UNDER (define subtype)

Controlling Access

- **Q:** How to allow user u_1 to see only (sid, name) of Student?
- **Q:** How to allow user u_1 to see only those students that take CS classes?

- **Q:** How to allow user u_1 modify a student record only if their age < 18 ?

- Authorization is one very important use of views.

Privileges and views

- **Q:** Assume user u is trying to create view V from R . User u has SELECT privilege on R (no grant option). Can u create such a view?
- **Q:** After V is created, can u grant SELECT on V to someone else?
- **Q:** If DBA revokes SELECT privilege on R from u , what should happen to V ?
- Notes on views and privileges
 - To create a view, a user needs the SELECT privilege on the base tables of the view
 - To grant a privilege on a view, the user has to have GRANT OPTION privilege on the base tables of the view
 - When the privileges needed for the base tables are revoked, the affected views are automatically dropped.